

GENERAL INSTRUCTION



Authors: Carefully check the page proofs (and coordinate with all authors); additional changes or updates WILL NOT be accepted after the article is published online/print in its final form. Please check author names and affiliations, funding, as well as the overall article for any errors prior to sending in your author proof corrections. Your article has been peer reviewed, accepted as final, and sent in to IEEE. No text changes have been made to the main part of the article as dictated by the editorial level of service for your publication.

Authors: We cannot accept new source files as corrections for your article. If possible, please annotate the PDF proof we have sent you with your corrections and upload it via the Author Gateway. Alternatively, you may send us your corrections in list format. You may also upload revised graphics via the Author Gateway.

AUTHOR QUERIES

- AQ1: Please note that there is a \$220 over-length charge per page beyond 10 pages for regular papers in this journal.
- AQ2: Please provide the ORCID for Paul Fearnhead in the byline.
- AQ3: Please check and confirm whether the funding information is correct as set.
- AQ4: Please specify the corresponding author's name.
- AQ5: Please provide the affiliation for all authors.
- AQ6: Operator symbols ($=$, $+$, and \times) have been repeated twice in the display equations throughout the article. Please check.
- AQ7: Please check whether the parenthesis style in this equation is correct.
- AQ8: Figures 5–9 are illegible. Please re-supply the figures with clear text.
- AQ9: Supplementary material is cited on page 8, but no material was provided with your source files. Please provide the missing files or rewrite the text to clarify what is being cited when you return your proof review.
- AQ10: Please provide the publisher location for [1], [22], and [33].
- AQ11: Please provide the volume number and issue number or month for [3], [17], and [35].
- AQ12: Please provide the volume number, issue number or month, and page range for [11].
- AQ13: Please provide the page range for [14].

A Log-Linear Nonparametric Online Changepoint Detection Algorithm Based on Functional Pruning

Gaetano Romano , Idris A. Eckley , and Paul Fearnhead

Abstract—Online changepoint detection aims to detect anomalies and changes in real-time in high-frequency data streams, sometimes with limited available computational resources. This is an important task that is rooted in many real-world applications, including and not limited to cybersecurity, medicine and astrophysics. While fast and efficient online algorithms have been recently introduced, these rely on parametric assumptions which are often violated in practical applications. Motivated by data streams from the telecommunications sector, we build a flexible nonparametric approach to detect a change in the distribution of a sequence. Our procedure, NP-FOCuS, builds a sequential likelihood ratio test for a change in a set of points of the empirical cumulative density function of our data. This is achieved by keeping track of the number of observations above or below those points. Thanks to functional pruning ideas, NP-FOCuS has a computational cost that is log-linear in the number of observations and is suitable for high-frequency data streams. In terms of detection power, NP-FOCuS is seen to outperform current nonparametric online changepoint techniques in a variety of settings. We demonstrate the utility of the procedure on both simulated and real data.

Index Terms—Changepoint, online, non-parametric, real-time analysis, anomaly detection, telecommunications.

I. INTRODUCTION

ONE of the contemporary challenges in time-series analysis is to detect changes in some measurable properties of a process. This task finds its roots in a plethora of applications spanning numerous fields including engineering [2], [22], neuroscience [24], genomics [30] and astrophysics [16], [18]. In the previous decade, we saw many offline changepoint procedures appearing in the statistical literature. A common approach for detecting a single change is to scan over all possible locations for a change, and apply a generalised likelihood ratio (GLR) test for a change, with evidence for a change being the maximum of the GLR test statistics [13]. The GLR procedure has been proven to be asymptotically optimal [5], as it searches over all possible parameters of the unknown pre and post-change distributions. This approach includes, as a special case, the popular CUSUM method for detecting a change in mean, and

can be extended to detect multiple changepoints by using binary segmentation methods [17], [26], [39] or by maximising a penalised likelihood [14]; see [11] for a recent review of this area.

In contrast to the offline setting, where the data is first collected and later analyzed *a posteriori*, one of the modern challenges is to detect a change within a data stream in real time. We find many applications in need of an online procedure. These include control of industrial processes [33], or monitoring of computer networks [41], social networks [6] and telecommunication devices [4].

Many challenges arise in analysing data online. An online procedure should be able to process observations in real-time, as quickly as they arrive, in order to avoid memory overflow and result in a delayed evaluation. This can be particularly valuable in settings where limited computational resources are available, or in the high-frequency domain.

As noted by [36], one way of performing an online analysis is to collect data in batches and analyse those through an offline algorithm. However, such an approach can be sensitive to the batch size. If this is too small, then we may be unable to detect small changes, while if it is large it can lead to delayed detection of bigger changes. Alternatively, observations can be processed on the go in a sequential fashion. At each new observation, a decision is made on whether a change has occurred, or not, based on the new data point and on past information.

The sequential CUSUM approach, and more generally, sequential GLR approaches, demonstrate excellent statistical properties [42], [46]. However, they can be computationally inefficient. A naive computation of the GLR test, in fact, involves, at time n , considering $\mathcal{O}(n)$ possible locations for a change. I.e., the algorithm has a computational cost per iteration that increases linearly, which is impracticable for online applications.

Recently, [34] presented the FOCuS procedure, a fast algorithm to perform the sequential CUSUM test, that decreases the computational complexity from $\mathcal{O}(n)$ to $\mathcal{O}(\log n)$ per iteration. This is to our knowledge the fastest way to solve the sequential CUSUM test exactly. For example, the expected cost per iteration of FOCuS at a given time, say at iteration one million, is roughly equal to the cost of iterating approximately 20 objects stored in memory to find the global maximum of our statistics – and thus it is suitable for high-frequency online applications.

As previously mentioned, the CUSUM and GLR tests in general rely on parametric assumptions that are often hard to

Manuscript received 3 February 2023; revised 30 June 2023 and 21 October 2023; accepted 6 December 2023. The work was supported in part by the EP-SRC and Lancaster University under Grants EP/N031938/1 and EP/R004935/1 and in part by the Mathematical Sciences RA under Grant EP/W522612/1. The associate editor coordinating the review of this manuscript and approving it for publication was Prof. Yao Xie. (Corresponding author: XXX.)

Digital Object Identifier 10.1109/TSP.2023.3343550




Fig. 1. Detection delays of Gaussian FOCuS (in red) and NP-FOCuS (in green) for 3 sequences of length 2000 with a change at $t = 1000$. In (a), a Gaussian change-in-mean example, with a jump size of 0.25. In (b), a change-in-variance from variance $\sigma^2 = 1$ to $\sigma^2 = 0.25$. Lastly, in (c), we find the same scenario in a with some of the observations shifted by a value of 3 at uniformly random sampled locations (blue dots).

meet in real-world applications. For instance, with an industrial collaborator, we found that certain data streams arising in the telecommunications sector do not follow common distributions, nor do they satisfy the usual parametric assumptions. In monitoring operational metrics from network devices, for example, we often deal with nonstandard distributions with multiple modes, outliers or heavy tails. In many cases, the nature of the change is unknown *a priori*, see for example the sequences from Fig. 8. In Section IV-B we present one example of a contemporary telecommunication application, where the aim is to monitor the operational performance of devices on an optical cable network. In those scenarios, a test for Gaussian change-in-mean would, in fact, be prone to either a missed detection, in the case of a change of a different nature, or to false positives, in case of misspecification of the underlying noise process.

To better outline the limitations of a parametric approach in online changepoint detection, we show a simple introductory example. Let us compare the Gaussian FOCuS from [34] with its non-parametric counterpart, NP-FOCuS, the methodology that will be introduced in this paper. In Fig. 1 we study both the statistics on 3 different simulated changepoint scenarios. Thresholds were tuned to achieve comparable average run lengths of 2000 observations under the null (the expected number of observations until we detect a change if a change is not present). This way, after placing a true change at 1000, we can have a fair comparison focusing on the detection alone. In Fig. 1(a), the simple Gaussian change-in-mean case, we notice how the Gaussian procedure achieves the fastest detection. This is because the procedure's assumptions hold. However in Fig. 1(b), the second example, the sequence now shows a change in variance. In this setting, the Gaussian change-in-mean has no power to detect the change, which is missed. Lastly, in

Fig. 1(c), we find the same series as in the first example, but some observations are shifted up by 3. Under this scenario, to achieve the same run length of 2000 observations under the null, the threshold for the Gaussian FOCuS procedure needs to be inflated. This results in a slower detection than its non-parametric counterpart.

Whilst there are numerous offline non-parametric approaches in the current literature (including and not limited to [7], [21], [28], [32], [47]), it is only more recently that there has been substantial interest on online approaches. This includes [38], who propose a control-chart approach that allows for the detection of changes in scale or location, or both, when the underlying distribution of the process is unknown, and [40] who propose a general framework for detecting changes in measurable properties of a process. More generally there are methods that choose a number of univariate summaries of data and monitor for a change in any of these summaries [25], [27], kernel-based methods [15], [23] or related methods that use information on distances between data points [6], [8], [12]. These approaches can be applied in multivariate settings, but their statistical efficiency depends on choosing appropriate choice of summary, kernel or distance that had power to detect the type of change that occurs. These methods are often not invariant to transformations of the data. By comparison, methods based on ranking information (e.g. [20]), are simple to apply for univariate data and are invariant to any monotone transformation of the data. This is also the case for methods that look for a change in the distribution function for univariate data, which is the approach we take in this paper.

One example of such a method is the NUNC algorithm [4], which tests for a change in the empirical cumulative distribution function (eCDF) within a rolling window. As with any other

moving window method, however, NUNC's performance is extremely dependent on the size of the window. For example, a window too small would end up missing changes of a smaller size, whilst a window too big might result in longer detection delays and an increased computational cost.

Our procedure, NP-FOCuS, tries to mitigate such issues by building a GLR test (extending ideas from [34]) for the non-parametric procedure NUNC of [4]. Our procedure maps the change-in-distribution problem into a Bernoulli GLR test. This is achieved by evaluating the eCDF at a fixed point and keeping track of the number of observations above or below that point. We then extend the FOCuS algorithm so that it can apply to the Bernoulli GLR test, and show theoretically that it has the same strong computational properties as the original FOCuS algorithm. We perform this test for a grid of quantiles, and then merge the test statistic values. In practice our final procedure monitors both the sum of the test statistics across the quantiles, which can detect small shifts in the distribution, and also the maximum of the test statistics, which can detect a larger change in just one part of the distribution. Our approach only assumes that the observations are *i.i.d.* However, we also demonstrate empirically that even if such a hypothesis is violated we still retain the strong performance over competitor methods.

The paper is structured as follows. In Section II we introduce the NP-FOCuS procedure, starting from related work, to the introduction of a novel functional pruning recursion. In Section III we provide some guarantees on the computational complexity of NP-FOCuS. In Section IV-A we evaluate empirically performances of NP-FOCuS, with a comparison to other methods in a simulation study. In Section IV-B we evaluate the method on some real-word data streams.

Open-source code, making the methods developed in this paper accessible, will be made available via an appropriate repository in due course.

II. A FUNCTIONAL REPRESENTATION FOR THE NONPARAMETRIC COST FUNCTION

A. Overview of Problem and Approach

We aim to build an online non-parametric changepoint detector by monitoring the empirical Cumulative Density Function (eCDF) of a stream of observations. This approach has been shown to work well in an offline setting by [21], [47].

Let $p \in \mathbb{R}$ and let $F_{1:n}(p)$ be the unknown CDF for a series of independent real-valued observations y_1, \dots, y_n , and let $\hat{F}_{1:n}(p)$ be its eCDF:

$$\hat{F}_{1:n}(p) = \frac{1}{n} \sum_{t=1}^n [\mathbb{I}(y_t \leq p)],$$

where \mathbb{I} is an indicator function. For an independent stream of observations, for a fixed p , the eCDF follows a binomial distribution:

$$n\hat{F}_{1:n}(p) \sim \text{Binom}(n, \theta),$$

with $\theta = F_{1:n}(p)$. A simple approach to detect a change in $\hat{F}_{1:n}(p)$ would be to record which data points are above p and

build a likelihood-ratio test for the Bernoulli data. That is, for a fixed p , we know that the likelihood for a segment is given by

$$\begin{aligned} \mathcal{L}(y_{\tau_1+1:\tau_2}; \theta, p) &= (\tau_2 - \tau_1) \times \\ &\times \left[\hat{F}_{\tau_1+1:\tau_2}(p) \log \theta + (1 - \hat{F}_{\tau_1+1:\tau_2}(p)) \log(1 - \theta) \right]. \end{aligned} \quad (1)$$

The GLR test for a change in parameter from θ_0 to θ_1 is

$$\max_{0 \leq \tau < n} [-\mathcal{L}(y_{\tau+1:n}; \theta_0, p) + \mathcal{L}(y_{\tau+1:n}; \theta_1, p)]. \quad (2)$$

By writing $x_t = \mathbb{I}(y_t \leq p)$, this can be rewritten as

$$\begin{aligned} &\mathcal{L}(y_{\tau+1:n}; \theta_1, p) - \mathcal{L}(y_{\tau+1:n}; \theta_0, p) \\ &= \sum_{t=\tau+1}^n [x_t \log \theta_1 + (1 - x_t) \log(1 - \theta_1)] \\ &\quad - \sum_{t=\tau+1}^n [x_t \log \theta_0 + (1 - x_t) \log(1 - \theta_0)] \end{aligned}$$

If both θ_0 and θ_1 are known, then the sequential procedure of [31] can be used to calculate the GLR test (2). In Section II-B we show how to extend the FOCuS algorithm of [34] to calculate the GLR test if either only θ_0 is known, or if neither θ_0 nor θ_1 is known. We will call this procedure Ber-FOCuS.

In practice, we would want to check changes in the entire eCDF rather than just at a single quantile value p . To achieve this we follow [47] and [21] who suggest summing up the test statistic across a grid of values p_1, \dots, p_M . Specifically, we approximate our eCDF on a fixed grid of M quantile values p_1, \dots, p_M computed over a probation period. To obtain NP-FOCuS we run the Ber-FOCuS routine independently for each quantile of our eCDF. To construct a global statistic we aggregate the quantile statistics through the sum and the maximum. In Section II-E we will formally describe the full procedure.

B. Detecting Change-in-Rate in a Bernoulli Process

We focus on detecting a change in the rate parameter θ in a univariate stream of a Bernoulli process. Let $x_n \in \{0, 1\} \sim \text{Ber}(\theta)$ be a realization of a Bernoulli random variable with parameter θ , with θ subject to a change. Assume the pre-change rate parameter, θ_0 , is known. For every observation we can obtain evidence for a post-change rate parameter θ_1 via the likelihood ratio test:

$$\begin{aligned} g(x_n, \theta_1) &= -2 \log \left(\frac{\theta_0^{x_n} (1 - \theta_0)^{1-x_n}}{\theta_1^{x_n} (1 - \theta_1)^{1-x_n}} \right) = \\ &= 2 \left[x_n \log \left(\frac{\theta_1}{\theta_0} \right) + (1 - x_n) \log \left(\frac{1 - \theta_1}{1 - \theta_0} \right) \right]. \end{aligned} \quad (3)$$

At time n we have observed x_1, \dots, x_n , we can then employ (4) and derive the test statistics:

$$\mathcal{Q}_{n, \theta_1} = \max_{0 \leq \tau < n} \sum_{t=\tau+1}^n g(x_t, \theta_1),$$

where we maximise for all possible time-points of the change, τ . Such a statistic is known as the sequential




Fig. 2. The functional cost $Q_n(\theta)$ at iteration n . Dotted lines illustrate (i) $\arg \max_{\theta} Q_n(\theta)$ and (ii) the $\max_{\theta} Q_n(\theta)$, the highest possible value of our test statistics.




Fig. 3. The components function of the functional cost $Q_n(\theta)$ at iteration n . The functions are plotted in different colours as dotted lines. The values for which they are optimal, e.g. they dominate all other lines and contribute to $Q_n(\theta)$, are solid. These curves are in line with those of the bottom right plot of in Fig. 4, illustrating the last step of our algorithm for $\theta > \theta_0 = 0.5$.

Page-CUSUM statistic [31], and can be solved efficiently in constant time per iteration through the recursion:

$$Q_{n,\theta_1} = \max \{0, Q_{n-1,\theta_1} + g(x_n, \theta_1)\}. \quad (5)$$

This statistic is straightforward to compute for a known value of θ_0 and a fixed value of θ_1 . As we do not know θ_1 , one would compute the statistics over a fixed grid of values for θ_1 and choose the maximum value of those to maximise the power of a detection. However, this would introduce a trade-off between power and computational complexity. Alternatively, [45] recently suggested a way of approximating the Page-CUSUM cost by using a plug-in value for $\theta_1 = \hat{\theta}$ into the equation above, that is calculated recursively based on only recent data point.

Our approach is to calculate the GLR test statistic exactly for both the case where θ_1 is unknown. The idea, based on [34], is to solve (5) simultaneously for all values of θ through the functional representation of the sequential test statistics $Q_n(\theta)$ for a post-change rate parameter θ . That is, we have $Q_0(\theta) = 0$ and:

$$Q_n(\theta) = \max \{0, Q_{n-1}(\theta) + g(x_n, \theta)\} \quad (6)$$

for $n \geq 1$. The value of our test statistic will be $Q_n = \max_{\theta} Q_n(\theta)$. In this optimization lies the major computational contribution of NP-FOCuS. At each iteration of the algorithm we explicitly compute and store the full functional representation of our cost $Q_n(\theta)$, which will be a piecewise function (see Fig. 2). By maximizing over all possible values of $\theta \in (0, 1)$, which is simple as we can maximise each piecewise part of the function and take the maximum of these maxima, we can find the highest possible value for the Page-CUSUM statistics, and avoid altogether any choice of window w or specify a post-change parameter. At a given iteration $Q_n(\theta)$ is obtained from a set of component functions of the form:

$$q^{(n,\tau)}(\theta) = a^{(n,\tau)} \log\left(\frac{\theta}{\theta_0}\right) + b^{(n,\tau)} \log\left(\frac{1-\theta}{1-\theta_0}\right), \quad (7)$$

with $a^{(n,\tau)} = \sum_{t=\tau}^n x_t$, $b^{(n,\tau)} = \sum_{t=\tau}^n (1-x_t)$. This curve represents the LR statistic for a change at τ .

Whilst there are n such possible curves, each one relative to a different candidate changepoint τ , in practice $Q_n(\theta)$ can be defined in term of a much smaller number of curves. In other words, for a large n , only a small proportion of these curves will contribute to the functional cost, *i.e.* will be the components of our piecewise function $Q_n(\theta)$, such that for some values of θ we have that $q^{(n,\tau)}(\theta) = Q_n(\theta)$. In our example, we plot in Fig. 3 the curves contributing to Fig. 2. If there are no values of θ for which a curve contributes to the optimal cost, we can drop it, a mechanism that is known as *pruning* that leads to the speed-ups in the algorithm. Therefore, Ber-FOCuS presents a way of iteratively and efficiently updating $Q_n(\theta)$, figuring out which are its components' functions and finding its global maximum to solve (6) exactly.

If k is the number of functions that contribute to $Q_n(\theta)$, we can represent $Q_n(\theta)$ on a computer efficiently as an ordered set of k triples $Q_n = \{q_i = (a_i^{(n,\tau_i)}, b_i^{(n,\tau_i)}, l_i^{(n,\tau_i)})\}$, $i = 1, \dots, k$. Each triple represents a curve $q^{(n,\tau_i)}(\theta)$, and contains the coefficients of $\log(\theta/\theta_0)$ and of $\log((1-\theta)/(1-\theta_0))$, and a value l_i which is the left-most value of θ for which the associated function is contributing to the optimal cost, *i.e.* $\theta \in [l_i^{(\tau_i,n)}, l_{i+1}^{(\tau_{i+1},n)})$, $q_i^{(n,\tau_i)}(\theta) = Q_n(\theta)$. In this interval our curve will be greater than all the other curves and the zero line.

For brevity, for the rest of this section, we will fix the iteration. We can simplify the notation by dropping the superscript (n, τ) , and hence we will denote our triples and functions simply as q_1, \dots, q_k . This is possible as n is fixed and τ is redundant information: by construction $\tau_i = n - (a_i + b_i)$ as we are dealing with binary data. We will assume that the triples are ordered in such a way that $\theta_0 = l_1 < \dots < l_k < l_{k+1} = 1$, so that for all $\theta \in [l_i, l_{i+1})$, $Q_n(\theta) = a_i \log(\theta/\theta_0) + b_i \log((1-\theta)/(1-\theta_0))$.

A description of the full Bernoulli FOCuS (Ber-FOCuS) procedure is given in Algorithm 1, with a graphical illustration of the procedure in Fig. 4. For each iteration, we have three main

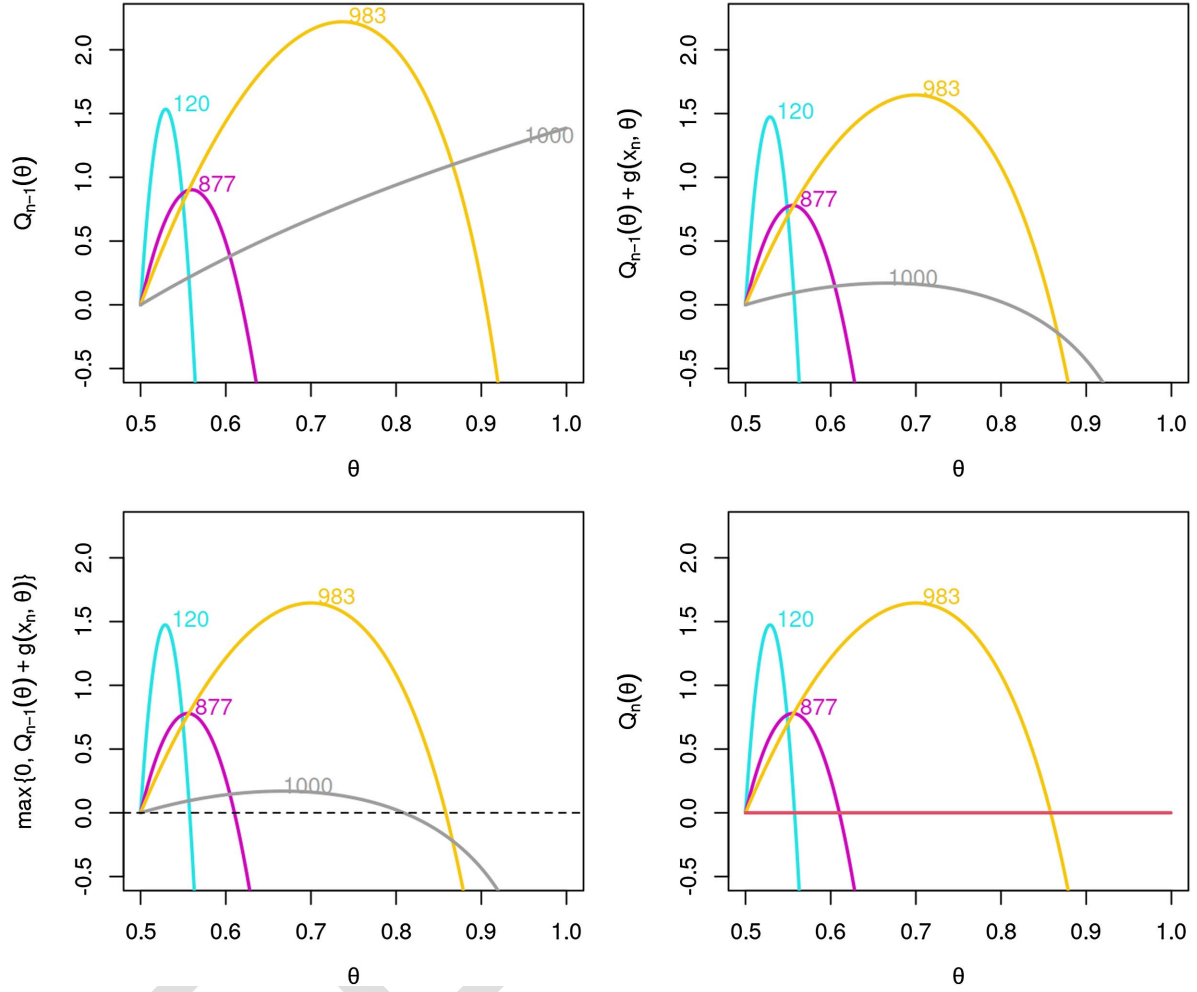


Fig. 4. One iteration of the NP-FOCuS procedure illustrated, for $\theta > \theta_0$. In this example, we assume $\theta_0 = 0.5$. Top-left: the cost function from the previous iteration $Q_{n-1}(\theta)$. The total cost is given by the maximum of 4 functions: we can see how each of those is optimal for some values of θ . Labels refer to the time τ at which each function was introduced. Top-right: the update step. By adding the new observation x_n , we update the coefficients $a_{i,\tau}$ or $b_{i,\tau}$ of each function. Although the shape of our function has changed, the points where quadratics intersect are unchanged. Bottom-left, the pruning step. Here we compare the functions with the zero line (dashed line). We find the lowest value of θ for which each function is optimal. We notice how for the grey function (1000) there is no such value of θ for which the function is optimal, as the intersection with its closest quadratic falls below the zero line. We can therefore safely prune the function. Bottom-right: $Q_n(\theta)$, after we add a new piece (in red).

Algorithm 1: Ber-FOCuS (one iteration)

Data: $x_n \in \{0, 1\}$ the data point at time n .

Input: $Q_{n-1}(\theta)$ the cost function from the previous iteration.

- 1 $\tilde{Q}(\theta) \leftarrow Q_{n-1}(\theta) + g(x_n, \theta)$
 - 2 $Q_n(\theta) \leftarrow \max\{0, \tilde{Q}(\theta)\}$; // see Algorithm 2
 - 3 $Q_n \leftarrow \max_{\theta} Q_n(\theta)$
 - 4 **return** $Q_n(\theta)$ for the next iteration, Q_n as the test statistic.
-

steps. Step 1, the update step, updates the values of the coefficients of stored functions according to (4). I.e., we update the count of positive or negative observations since the introduction of a curve. This is a simple increment of either coefficient a_i or b_i . Step 2, which will be detailed in Section II-C, is the

functional pruning step and focuses on identifying which pieces are no longer optimal (and will never be in the future). Finally, the optimization step, where we maximise the total cost $Q_n(\theta)$ over θ , to get our test statistics. This is equivalent to taking the maximum of all maximums of our functions. As we will see in Section III to reconstruct the optimal cost we need only to store a small subset of candidate changepoints: at the n^{th} iteration we expect to store on average $k \approx \log(n)$ functions.

C. Step 2: Efficient Pruning of the Bernoulli Cost

Functional pruning consists of restricting the set of values to consider as candidate changepoints on account of the information they carry up to time n . That is, at each iteration we can stop considering functions introduced at times if they will never be optimal for our cost $Q_n(\theta)$. To achieve this, [34] propose an efficient pruning rule, we can adapt to our setting. As in [34], we propose to update $Q_n(\theta)$ separately for $\theta > \theta_0$ and $\theta < \theta_0$. We

will report the update rule for $\theta > \theta_0$ alone as, by symmetry, the update for $\theta < \theta_0$ can be found by applying the same algorithm with the data inverted, i.e. replacing x_t with $(1 - x_t)$.

If we consider Ber-FOCuS algorithm of Algorithm 1, then at the start of an iteration we will have a set of triples. The first step, the update step, only affects the a_i and b_i coefficients. The l_i components, which specify the intervals of θ that each curve is optimal, are unchanged. This follows as the differences between the curves are unaffected by the update, as each curve is changed by adding the same function $g(x_n, \theta)$, and it is the difference between curves that determines where one is greater than all others.

The pruning of curves only occurs in step 2, and is as a result of taking the maximum of $\tilde{Q}(\theta)$ and the zero line. The idea is that we need to find the values of θ for which the zero line is optimal. Define $\text{root}(q_i)$ to be the largest value of $\theta > \theta_0$ such that $q_i(\theta) = 0$, if such a value exists. It is straightforward to show that, for a given function q_i that contributes to $\tilde{Q}(\theta)$ for $\theta > \theta_0$, if $\text{root}(q_i) < 1$ exists, then the zero-line is better on the interval:

$$[\text{root}(q_i), 1).$$

Considering all functions, this implies that if $\max_{\tau} \text{root}(q_i) < 1$ then the zero line is globally optimal on $[l^*, 1)$, where

$$l^* = \max_{\tau} \text{root}(q_i).$$

Therefore, the triple defining the zero line will be given by $(0, 0, l^*)$. Given that the quadratics are ordered, it means that any quadratic with $l_i < l^*$ can therefore be removed. Finally, the ordering of the quadratics means that $l_i < l_j$ if and only if $i < j$. For this reason, we can start checking from the last quadratic q_k and stop as soon as the pruning condition is not met. The pruning procedure is summarised in Algorithm 2. Following the update, for pruning, given that at any point in time $l_i = \text{root}(q_i - q_{i-1}) \leq l^*$ and that $q_i(\theta_0) = 0$, we know that whether $q_i(l_i) < 0$ there will be no such value of θ for which q_i will be optimal. In fact, we know that q_i will be greater than q_{i-1} only starting from l_i , and if $q_i(l_i) < 0$ falls below the zero line, then $q_{i-1}(\theta) < q_i(\theta) < 0 \forall \theta \in (l_i, 1)$ and therefore q_i will be never be optimal and can be safely pruned. The bottom-left plot of Fig. 4 can be of aid in understanding the pruning and the following theorems.

We therefore learn that the algebraically ideal version of the algorithm relies on computing the value $\text{root}(q_k)$ in step 9 at each iteration. However, given the shape of our cost function, we do not have a closed form for computing this root. A naive implementation would thus require obtaining such a value numerically.

Fortunately, there is an alternative way of performing the same pruning exactly that avoids computing the root value. As we will learn in the following Theorem and Lemma, this will be achieved by studying the link between the value of the roots of our pieces and their argmax.

Theorem 1: Assume q_i and q_j are two curves that contribute to $Q_n(\theta)$ for $\theta > \theta_0$, with $i > j$. Then, $\text{root}(q_i) < \text{root}(q_j)$ if and only if $a_i/(a_i + b_i) < a_j/(a_j + b_j)$.

Algorithm 2: Algorithm for $\max\{0, \tilde{Q}(\theta)\}$ for $\theta > \theta_0$

Input: \tilde{Q} , an ordered set of triples

$$\{q_i = (a_i, b_i, l_i) \mid i = 1, \dots, k\}$$

```

1  $i \leftarrow k$ ;
2 while  $q_i(l_i) \leq 0$  and  $i \geq 1$  do
3    $i \leftarrow i - 1$ ;
4 end
5 if  $i \neq k$  then
6    $Q \leftarrow Q \setminus \{q_{i+1}, \dots, q_k\}$ ; // drop pieces
7    $k \leftarrow i$ 
8 end
9 if  $k = 0$  then
10   $q_{k+1} \leftarrow (0, 0, \theta_0)$ ; // add first piece
11 else
12   $q_{k+1} \leftarrow (0, 0, \text{root}(q_k))$ ; // add new piece
13 end
14  $Q \leftarrow \{Q, q_{k+1}\}$ 
15 return  $Q$ 

```

Proof: We begin with some properties of the curves. A curve q_i has $\arg \max q_i(\theta) = a_i/(a_i + b_i)$. By construction, if $i > j$ then $a_i \leq a_j$ and $b_i \leq b_j$. Finally $q_i(\theta_0) = q_j(\theta_0) = 0$.

We now link the arg max of our functions to their roots. First we can rescale each curve without changing the location of its maximum of its root. For curve q_i define $\tilde{a}_i = a_i/(a_i + b_i)$, $1 - \tilde{a}_i = b_i/(a_i + b_i)$ and write

$$\tilde{q}_i(\theta) = \tilde{a}_i \log\left(\frac{\theta}{\theta_0}\right) + (1 - \tilde{a}_i) \log\left(\frac{1 - \theta}{1 - \theta_0}\right). \quad (8)$$

We then have that

$$\text{root}(\tilde{q}_i) \geq \text{root}(\tilde{q}_j) \iff \tilde{a}_i \geq \tilde{a}_j, \quad (9)$$

This follows by noting that $\tilde{q}_i(\theta)$ is unimodal as $\frac{\partial \tilde{q}_i(\theta)}{\partial \theta}$ has one zero, and that $\forall \theta \in [0, 1]$, $\frac{\partial \tilde{q}_i(\theta)}{\partial \theta} \geq \frac{\partial \tilde{q}_j(\theta)}{\partial \theta} \iff \tilde{a}_i \geq \tilde{a}_j$.

As q_i has the same root as \tilde{q}_i , the result follows immediately. \square

We will use this result to simplify the pruning procedure in the following Lemma.

Lemma 2: The condition for pruning $q_i(l_i) \leq 0$ at step 2 of Algorithm 2 is implied by $a_i/(a_i + b_i) < a_{i-1}/(a_{i-1} + b_{i-1})$ for $i > 1$ and $a_i/(a_i + b_i) < \theta_0$ for $i = 1$.

Proof: To link the pruning condition and the results from Theorem 1 we need to show that for $i > 1$ whether $q_i(l_i) \leq 0$ we find $\text{root}(q_i) < \text{root}(q_{i-1})$. We start by noting that l_i is the root of the function $q^*(\cdot) = q_i(\cdot) - q_{i-1}(\cdot)$, i.e. $l_i = \text{root}(q^*)$, and that $q^*(\theta) \geq 0$ on $[\theta_0, l_i]$.

Then, if $q_i(l_i) \leq 0$ we have

$$q_{i-1}(\theta) = q_i(\theta) + q^*(\theta) > q_i(\theta) \text{ for } \theta \in (\theta_0, l_i),$$

and as $\text{root}(q_{i-1}) \in (\theta_0, l_i)$, it follows that

$$q_i(\text{root}(q_{i-1})) < q_{i-1}(\text{root}(q_{i-1})) = 0.$$

By concavity of q_i then $\text{root}(q_i) < \text{root}(q_{i-1})$. The argument for $i = 1$ follows simply by noting that in this case $l_1 = \text{root}(q_1) < \theta_0 \iff a_1/(a_1 + b_1) < \theta_0$. \square

Hence we can swap the pruning condition involving the root-finding with the simpler condition: $a_i/(a_i + b_i) < \max\{\theta_0, a_i/(a_i + b_i)\}$. Given that this novel pruning only involves the a_i, b_i values, it is no longer necessary to find and store the l_i value at each iteration. This result relies, in part, on the functions, $q_i(\cdot)$, being unimodal. Similar results apply for other functions, and these have been used to generalise the FOCuS algorithm to other one-parameter exponential family models (see [43]).

D. Extension to the θ_0 Unknown Case

We next extend our recursion to the case where both θ_0 and θ_1 are unknown. Assume we are observing a stream of Bernoulli random variables x_1, \dots, x_n distributed as a $Ber(\theta_0)$ under the null and as a $Ber(\theta_1)$ under the alternative. The likelihood ratio test statistic is

$$\mathcal{Q}_n = -\max_{\theta \in \mathbb{R}} \sum_{t=1}^n h(x_t, \theta) + \quad (10)$$

$$+ \max_{\tau \in \{1, \dots, n-1\}} \left\{ \sum_{t=1}^{\tau} h(x_t, \theta_0) + \sum_{t=\tau+1}^n h(x_t, \theta_1) \right\}, \quad (11)$$

where

$$h(x_t, \theta) = x_t \log \theta + (1 - x_t) \log(1 - \theta).$$

Solving this directly for all possible values of τ , via directly storing the partial sums, will result in a procedure that has computational complexity of $\mathcal{O}(n)$ per iteration and $\mathcal{O}(n)$ in memory. We employ a functional pruning approach and solve (11) exactly through the following recursion.

Proposition 3: Let $Q_0(\theta) = 0$ and for $n = 1, 2, \dots$ let Q_n be defined by the recursion:

$$Q_n(\theta) = \max \left\{ \max_{\theta} \sum_{t=1}^n h(x_t, \theta), Q_{n-1}(\theta) + h(x_n, \theta) \right\}.$$

Then, at time n , $\mathcal{Q}_n = \max_{\theta} \sum_{t=1}^n h(x_t, \theta) - \max_{\theta} Q_n(\theta)$.

The proof to Proposition 3 is found in Section A of Appendix.

This recursion can be solved using the same ideas as from [35]. As before $Q_n(\theta)$ will be the maximum of a set of curves. Each curve will be of the form

$$a \log(\theta) + b \log(1 - \theta) + c,$$

where the coefficients will depend on the changepoint location associated with that curve. Importantly the set of changepoint locations that we need to store curves correspond to locations whose curves would be stored for the θ_0 known case for some θ_0 . The only difference that θ_0 has on the changepoint locations whose curves are kept is that when considering positive changes we only need consider $\theta_1 > \theta_0$, and for negative changes we need only consider $\theta_1 < \theta_0$. As we need to keep curves that are optimal for some value of θ_0 this means that for positive changes we need to prune the same curves as for the θ_0 known case by for $\theta_0 \rightarrow 0$. For negative changes we prune as for the θ_0 known case but for $\theta_0 \rightarrow 1$.

Algorithm 3: NP-FOCuS (one iteration)

Data: $y_n \in \mathbb{R}$ the data point at time n .

Input: $\xi^{sum}, \xi^{max}, \{p_1, \dots, p_M\}, \{Q_{n-1}^1(\theta), \dots, Q_{n-1}^M(\theta)\}$.

```

1 for  $m = 1, \dots, M$  do
2    $Q_n^m(\theta), \mathcal{Q}_n^m \leftarrow \text{Ber-FOCuS}(x_n = \mathbb{I}(y_n < p_m); Q_{n-1}^m(\theta));$  // Algorithm 1
3 end
4  $\mathcal{S}^{sum} \leftarrow \sum_{m=1}^M \mathcal{Q}_n^m;$ 
5  $\mathcal{S}^{max} \leftarrow \max_{m \in \{1, \dots, M\}} \mathcal{Q}_n^m;$ 
6 if  $\mathcal{S}^{sum} \geq \xi^{sum}$  or  $\mathcal{S}^{max} \geq \xi^{max}$  then
7   return  $n$  as a stopping point.
8 end
9 return  $\{Q_n^1(\theta), \dots, Q_n^M(\theta)\}$  for the next iteration.
```

The algorithm to implement Ber-FOCuS in the θ_0 unknown for $\theta_1 > \theta_0$ is almost identical to Algorithm 2 except we now store a 4-tuple (a_i, b_i, c_i, l_i) for each curve, so in step 9 the new curve is described by the 4-tuple $(0, 0, c, \text{root}(q_k))$ with $c = \max_{\theta} \sum_{t=1}^n h(x_t, \theta)$. As before, in practice we can implement the resulting algorithm without storing l_i .

E. Aggregating the Bernoulli Traces and NP-FOCuS

Having covered how to independently monitor the M Bernoulli processes for all quantiles, we describe how to obtain a global statistic for NP-FOCuS.

Our approach is to consider two aggregation functions. One is to take the maximum of the statistics, and the other is to take the sum, as in e.g., [29].

NP-FOCuS: Let $Q_n^1(\theta), \dots, Q_n^M(\theta)$ be the costs at time n for the M Bernoulli sequences $\mathbb{I}(y_n \leq p) p \in \{p_1, \dots, p_M\}$, and let the Ber-FOCuS statistics:

$$\mathcal{Q}_n^m = \max_{\theta} Q_n^m(\theta)$$

Then, we will detect a changepoint at time n whether:

$$\sum_{m=1}^M \mathcal{Q}_n^m \geq \xi^{sum} \text{ or } \max_{m \in \{1, \dots, M\}} \mathcal{Q}_n^m \geq \xi^{max},$$

with $\xi^{sum}, \xi^{max} \in \mathbb{R}$. A formal description of the NP-FOCuS algorithm is reported in Algorithm 3.

One drawback of having to monitor two streams from two different aggregators comes with the procedure requiring two thresholds. However, there is an advantage in monitoring both streams simultaneously, as there are change scenarios where we expect one to perform better than the other. The sum statistic should have greater power to detect small shifts in the distribution that affect all or many quantiles. The maximum statistic will have greater power to detect larger shifts that affect e.g. just the tail of the distribution. By setting either ξ^{sum} or ξ^{max} to infinity, our method would result to considering a test just based on the maximum or the sum.




Fig. 5. Average Detection Delay in function of the number of quantiles M . Vertical bars represent the standard deviation across the various replicates.

472 F. Tuning Strategy and Choice of Quantiles

473 We now describe how to tune the initial parameters of the
 474 NP-FOCuS procedure. Quantile values can be computed either
 475 on training observations or on a probation period, in case of no
 476 training data. Rather than taking evenly spaced quantiles, we
 477 build geometrically spaced quantiles following the approach of
 478 [21]. That is, for fixed M , we take p_1, \dots, p_M in such a way
 479 that p_m is the empirical quantile with probability

$$\left\{ 1 + (2n - 1) \exp \left[-\frac{(2m - 1)}{M} \log(2n - 1) \right] \right\}^{-1}.$$

480 This is to give more importance to values in the tail of
 481 the distribution of the data. In case of a change in the tails
 482 of a distribution, for instance, this would allow for a quicker
 483 detection as the traces of the more-extreme quantiles are more
 484 likely to pass the threshold for the max statistic.

485 As we will find out from Section III, the choice of M affects
 486 the computational complexity of the procedure. Even if a higher
 487 number of quantiles should lead to better statistical power, in
 488 practice there is not much of a gain in picking values of M
 489 greater than 15. This can be checked from the elbow plot in
 490 Fig. 5, where we measure the average detection delay as a
 491 function of M . Simulations were performed as described in
 492 Section IV-A.

493 Finally, we tune the thresholds ξ^{sum} or ξ^{max} via a Monte
 494 Carlo approach. This is common practice in the literature when
 495 using fixed thresholds. In particular, we follow the approach of
 496 [9, Section 4.1]. We generate, say, T many sequences as long
 497 as the desired average run length N under the null. This can
 498 be achieved either by simulating the process directly or boot-
 499 strapping from training data. We then compute both the sum
 500 and the max statistics on all the T sequences and pick the max
 501 on each. The idea is that the stopping times S^{sum}, S^{max} , the
 502 times in which respectively the sum or the max statistics go over
 503 the thresholds ξ^{sum}, ξ^{max} , are approximately exponentially
 504 distributed. This means that we want to choose the thresholds
 505 so that the probability of stopping after time N , if there is no
 506 change, is $\exp(-1)$.

507 Based on our simulations we pick the estimated thresholds
 508 $\tilde{\xi}^{sum}, \tilde{\xi}^{max}$ to be the smallest values such that the proportion

of times we do not detect a change by time N , based on the
 sum and max statistics respectively, is less than $\exp(-1)$. This
 approach would give thresholds with the approximate average
 run length we require if we were using only the max or only
 the sum statistic. As we are using both, we now fix the ratio of
 $\tilde{\xi}^{sum}$ to $\tilde{\xi}^{max}$ but use the same procedure to scale their values
 to get a composite test with the correct average run length.

III. COMPUTATIONAL COMPLEXITY OF THE NP-FOCuS PROCEDURE

The computational and memory complexity per iteration of
 Algorithm 3 is M times the cost of running Ber-FOCuS. The
 computational cost of Ber-FOCuS consists of the pruning step
 and the update and maximisation steps.

The pruning step has a computational cost that is, on average
 $O(1)$ per iteration. This is because at each iteration we will
 consider some number, m say of curves – and this will result in
 at least $m - 1$ curves being removed. As we only add one curve
 per iteration, and a curve can only be removed once, this limits
 the average number of curves to be considered per iteration to
 be at most 2.

The other steps have a computational cost that is proportional
 to the number of curves that are kept. Furthermore, empirically,
 these step dominate the computational cost of the algorithm. We
 can bound the expected cost of these steps with the following
 result which bounds the number of curves that are stored.

Theorem 4: Let x_1, \dots, x_n, \dots be a realization of an in-
 dependent Bernoulli process centred on θ_0 . Let the number of
 functions stored by Ber-FOCuS for $\theta > \theta_0$ at iteration n be
 $\#\mathcal{I}_{1:n}$. Then if there is no change prior to n

$$E(\#\mathcal{I}_{1:n}) \leq (\log(n) + 1),$$

while if there is one change prior to n then

$$E(\#\mathcal{I}_{1:n}) \leq 2(\log(n/2) + 1).$$

The proof of this theorem is found in Section A of Appendix.
 By symmetry, the theorem extends to values of $\theta < \theta_0$. From the
 theorem, we learn that, at each iteration, we expect to check
 $k = \log(n) + 1$ curves in case a change has not occurred, and
 $k = 2 \log(n/2) + 1$ curves in the case where a change has been
 occurred but is undetected.

IV. EMPIRICAL EVALUATION OF THE METHOD

A. Simulation Study

We performed a simulation study to assess performances
 of NP-FOCuS and other online changepoint procedures in a
 variety of scenarios, illustrated in Fig. 6. Those scenarios were
 chosen to benchmark procedures over a range of different chal-
 lenges in online changepoint detection. Specifically, present
 scenarios for Cauchy change-in-scale, Gaussian change-in-
 mean, change-in-mode in a mixture of two Gaussian distribu-
 tions, change-in-mean in an Ornstein-Uhlenbeck process, decay
 in a sinusoidal process with random noise, and a change-in-tails
 scenario – we cover in details how to generate these in the Sup-
 plementary Materials, Section B. All the scenarios show *i.i.d.*

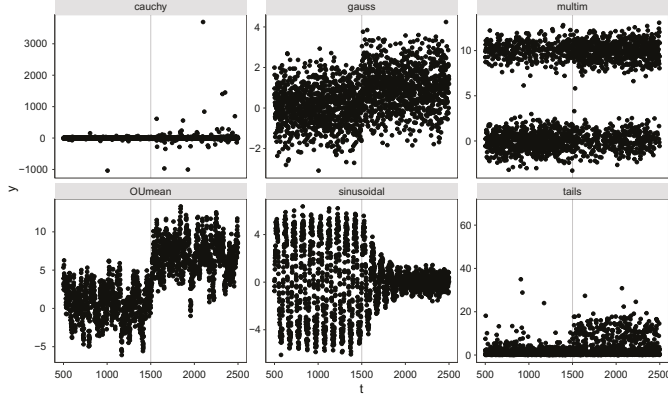


Fig. 6. Example sequences from the six different scenarios considered for the simulation study. Solid grey lines demarcate a changepoint.

sequences with the exception of the OU process and the Sinusoidal process. These two were added to account for changes in scale and location in presence of strong temporal dependency, a well-known challenge in changepoint detection which is present in many real-world applications (see [10], [19], [34]).

We compare NP-FOCuS with the online non-parametric procedure NUNC [4], the method from [37], the NEWMA procedure with RFF (see Section IV of [25] NEWMA and the recent online kernel CUSUM procedure from [44] with Gaussian Kernel (that we call Wei-CUSUM). As with other methods, Wei-CUSUM procedure assumes independent data. However, whereas the other methods can be applied to non-independent data by increasing the threshold for detecting a change, this is impossible for such a procedure. In fact, in line with other Kernel based online statistics, such as KCUSUM from [15], if the data is serially dependent then the contribution to the statistic for each new data point will on average be positive even if there is no change – and thus the method will be prone to false positives regardless of how large the threshold for the test is chosen. Thus we exclude Wei-CUSUM from the sinusoidal and OU scenarios. Lastly we add to the comparison the FOCuS procedure for Gaussian change-in-mean from [34]. A robust bi-weight loss was employed in this case to account for the presence of outliers in some scenarios.

Each experiment consisted of 100 replicates. Initial parameters were tuned according to Section II-F. Thresholds were selected to achieve an average run length of 10,000 observations. Other parameters – such as quantiles, the value of the bi-weight loss parameter or the NEWMA bandwidth – were obtained over a probation period consisting of the first 100 observations (in Section C of Appendix we study the effect of the length of the probation on NP-FOCuS performances). The NEWMA method involves monitoring the difference of two exponentially weighted moving averages of features of the data, and it can take a substantial number of observations for this difference to stabilise. To account for this we only start monitoring the NEWMA statistic after a burn-in period of 1000 observations to avoid false-positives whilst its statistics stabilise. For Wei-CUSUM, we gave oracle knowledge of the pre-change distribution, and tuned the scale parameter of the Gaussian kernel based on the variance of the pre-change distribution.

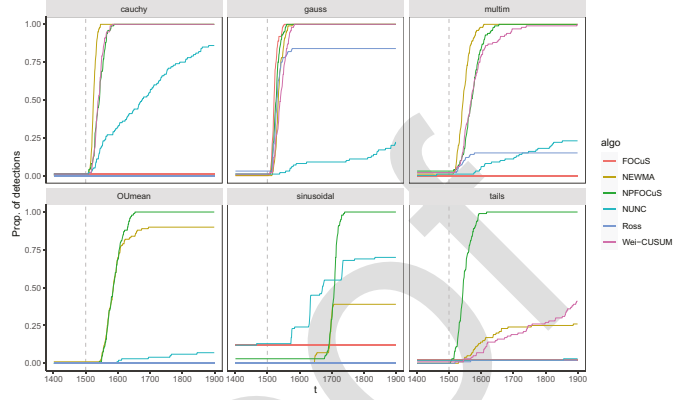


Fig. 7. Proportions of changes detected within t observations following the change in six different change scenarios. The change is denoted by the vertical dotted line at $t = 1000$.

We measure performances in terms of detection delay on sequences generated to have a change at $\tau = 1500$. In Fig. 7, we report the proportions of experiments where a change was detected by time step t . Prior to the change, the lines show the proportion of false positives, while following the change, they show the number of true detections by a given time step: the perfect online procedure would achieve a detection within 1 observation from the change in all sequences. Additionally, we report results in terms of average detection delay and false positive rate in Table I(a) and I(b).

We learn that, overall NP-FOCuS shows good performances in terms of statistical power. In the Gaussian case, unsurprisingly, the simple FOCuS with Gaussian loss has best performances overall, immediately followed by NP-FOCuS. Additionally, we found that NEWMA achieves faster detections both in the multimodal scenario, and in the Cauchy scenario, immediately followed by NP-FOCuS and Wei-CUSUM with similar performances. However, upon further testing, we noticed that the performance of NEWMA tends to degrade if the change occurs earlier, even if we use a shorter burn-in period (see Section C of Appendix for further details). Lastly, Kernel based procedures such as Wei-CUSUM, in addition to requiring the ability to sample from the pre-change distribution, are sensitive to the choice of the kernel.

We find NP-FOCuS to be more robust to strong dependence in the signal than other procedures. The most challenging scenario for NP-FOCuS is the sinusoidal scenario, where we start to consistently estimate the change with a delay of 200 observations. This, as it can be seen from Fig. 6, corresponds to the point where there is a clear difference in scale. In this scenario, for comparison, other methods are either prone to false positives, or missed detections.

B. Monitoring Power Attenuation on Optical Lines

We now evaluate NP-FOCuS on a real-world application, that of monitoring power attenuation on an optical communication line, a metric that is associated with the operational performance of the line. We present in Fig. 8 four examples of such a metric. Under normal circumstances, the power attenuation should show a stationary behaviour over time, as seen in Fig. 8(a).

TABLE I
AVERAGE DETECTION DELAY AND FALSE POSITIVE RATE ACROSS THE VARIOUS ALGORITHMS FOR ALL SCENARIOS OF OUR SIMULATION STUDY

scenario	FOCuS	Wei-CUSUM	NEWMA	NPFOCuS	NUNC	Ross
OUMean	>1000		225.92	87.99	995.35	>1000
cauchy	>1000	35.03	24.52	33.98	242.72	>1000
gauss	22.08	27.2	35.71	22.26	680.64	238.87
multim	>1000	65.44	42.54	44.86	943.52	>1000
sinusoidal	>1000		988.15	165.8	423.91	>1000
tails	>1000	745.03	>1000	46.97	>1000	>1000

(a) Average detection delay. In bold, the best result by row

scenario	FOCuS	Wei-CUSUM	NEWMA	NPFOCuS	NUNC	Ross
OUMean	0.00		0.01	0.00	0.00	0.00
cauchy	0.01	0.01	0.00	0.01	0.01	0.00
gauss	0.00	0.01	0.00	0.01	0.01	0.03
multim	0.00	0.02	0.01	0.03	0.01	0.02
sinusoidal	0.12		0.00	0.03	0.13	0.00
tails	0.02	0.01	0.00	0.00	0.02	0.00

(b) False positive rate




Fig. 8. Four examples of the measured power attenuation over time. Red segments at the bottom of the sequences point to detection times of FOCuS for Gaussian change-in-mean (red), the presented procedure NP-FOCuS (green) and finally NUNC (blue). Respectively, we show, in (a), a sequence with no change, in (b), a sequence with a shift in some modes, in (c), a sequence with changes in the behaviours of the tails, and finally in (d) an abrupt change-in-location.

Any change in distribution might therefore be of interest to the engineers and be an object of further investigation. Efficient techniques are required as there is a need of monitoring a large number of instances over time.

Each time-series has recordings over 80 days. The time-series were first classified by a domain expert into stationary series, and series with a change that would want to be flagged to an engineer. Thresholds were tuned over 85 non-changing sequences to achieve a false positive rate of 0.01, as described in Section II-F, with quantiles being trained over a probation period of 2 weeks in each sequence.

We draw a comparison of NP-FOCuS stopping times with both Gaussian FOCuS and NUNC. Threshold and other tuning parameters were adjusted on the same NP-FOCuS training instances to achieve comparable false positive rates. All 3 algorithms do not report a detection on the example without a change. In the sequences with a change, NP-FOCuS reported faster detection times. Overall, Gaussian FOCuS showed the slowest detection delay, even for the example where the change appears to be a change-in-location. This is attributable to the fact that the training data is heavy-tailed, and as mentioned in the example in the introduction this leads to slower detections

caused by over-inflated thresholds. Comparing the nonparametric online procedures, NP-FOCuS and NUNC, they show similar detection times in the b and d sequences, with NP-FOCuS resulting in slightly faster detections. However, in line with the previous numerical study, NUNC struggles in detecting changes in the tails of a distribution.

V. REMARKS AND CONCLUSION

We presented a nonparametric approach for online change-point detection. The procedure keeps track of the eCDF of a data stream over a set of fixed quantiles. Keeping track of which observations are above or below each quantile maps the problem into a Bernoulli change-in-rate problem. We derived a functional pruning sequential LR test to perform such a task efficiently in almost linear time.

While we have provided theoretical guarantees about the computational aspects of our procedure, we have not given any results on the average run length or the expected detection delay. Obtaining tight results is challenging for our approach due to additional complexity of allowing for any post change (and potentially pre-change) parameter, and due to combining of information across different quantiles. For theoretical results for the related off-line changepoint procedure, including showing that in the limit where we increase the number of quantiles, we are able to have power to detect any change in distribution, see [47].

The only assumption required by the NP-FOCuS, that of i.i.d. observations, can lead to a drop in performances when dealing with time-dependent sequences. In presence of strong dependence, to achieve comparable run lengths to the i.i.d. case, thresholds are to be inflated, resulting in a slower detection delay. An approach that can account for the time dependence, such as those of [10], [35], could be an avenue for further research. Relying on fixed quantiles also means that the procedure needs some training observations to understand the range of the data. This can be a problem in particular when restarting the procedure. If the quantiles are not known, in fact, those are needed to be re-estimated at every detected change. This means that no detection will be made for the entire duration of the probation period. One simple solution would be to store a rolling window of observations of the length of the probation period. Then, as soon as a detection is made, the procedure could be restarted with quantiles computed over those stored observations.

APPENDIX

A. Proofs

Proof of Proposition 3

For $n = 0$ the recursion follows trivially. For $n > 0$, at each iteration, we can either have

$$\begin{aligned} \mathcal{Q}_n &= \max_{\theta} \sum_{t=1}^n h(x_t, \theta) - \max_{\theta} \mathcal{Q}_n(\theta) = \\ &= \max_{\theta} \sum_{t=1}^n h(x_t, \theta) - \max_{\theta} \sum_{t=1}^n h(x_t, \theta), \end{aligned}$$

if there is no change up to time n , or for a change at time $\tau \in 1, \dots, n-1$:

$$\begin{aligned} \mathcal{Q}_n &= \max_{\theta} \sum_{t=1}^n h(x_t, \theta) - \max_{\theta} \mathcal{Q}_n(\theta) = \\ &= \max_{\theta} \sum_{t=1}^n h(x_t, \theta) - \max_{\theta} [\mathcal{Q}_{n-1}(\theta) + h(x_n, \theta)] = \\ &= \max_{\theta} \sum_{t=1}^n h(x_t, \theta) \\ &\quad - \max_{\theta} \left[\max_{\theta_0} \sum_{t=1}^{\tau} h(x_t, \theta_0) + \sum_{t=\tau+1}^{n-1} h(x_t, \theta) + h(x_n, \theta) \right]. \end{aligned}$$

□

Proof of Theorem 4

The proof follows almost exactly the one derived for Theorem 3 in [34]. While they assume a piecewise constant process plus i.i.d. noise from a continuous distribution, we denote how this assumption can be relaxed to include any exchangeable real-valued random process with an additive cost function, such as our Bernoulli process. Their proof relies on the following lemmas.

Lemma 5: For $i \leq j \leq k$

$$\mathcal{I}_{i:k} \subseteq \mathcal{I}_{i:j} \cup \mathcal{I}_{j+1:k} \quad (12)$$

Lemma 6: The set of τ in $\mathcal{I}_{i:j}$ are the extreme points of the largest convex minorant of the sequence $S_{i:j} = \left\{ \sum_{t=i}^i x_t, \dots, \sum_{t=j}^j x_t \right\}$.

We notice that if we write the cost of a segment with a change at $\tau \in \mathbb{N}$ as

$$\begin{aligned} q_{i:j,\tau}(\theta_0, \theta_1) &= \sum_{t=i}^{\tau} [x_t \log \theta_0 + (1 - x_t) \log(1 - \theta_0)] \\ &\quad + \sum_{t=\tau+1}^j [x_t \log \theta_1 + (1 - x_t) \log(1 - \theta_1)], \end{aligned}$$

then for any τ, τ' :

$$\begin{aligned} q_{i:j,\tau}(\theta_0, \theta_1) - q_{i:j,\tau'}(\theta_0, \theta_1) &= \\ &= \sum_{t=\tau}^{\tau'} \left[x_t \log \left(\frac{\theta_1}{\theta_0} \right) + (1 - x_t) \log \left(\frac{1 - \theta_1}{1 - \theta_0} \right) \right] \end{aligned}$$

this does not depend on i and j . This allows us to extend Lemmas 5 and 6 to the Bernoulli case.

And as both [3] and [1] assume an exchangeable real-valued random variable, we can adapt Lemma 7 of [34] to the Bernoulli case:

Lemma 7: Assuming the x_t follows an exchangeable real-valued process $\forall t \in i : j$, then $E(\#\mathcal{I}_{i:j}) = \sum_{t=1}^{j-i-1} 1/(t+1)$.

Having shown that the three lemmas apply to the Bernoulli case we can simply follow the proof from [34]. □

B. Detailed Description of Simulation Scenarios

In this section we describe the simulation scenarios, detailing the sampling for respectively the pre-change and post-change distributions. The Cauchy change-in-scale was obtained by sampling $y_{1:\tau}$ from a Cauchy with location 0 and scale 1, i.e., $y_{1:\tau} \sim \text{Cauchy}(0, 1)$ and for post change $y_{\tau+1:n} \sim \text{Cauchy}(0, 5)$. Following similar notation, the Gaussian change-in-mean was obtained by $y_{1:\tau} \sim N(0, 1)$ and $y_{\tau+1:n} \sim N(1, 1)$. The multimodal scenario was obtained by sampling $y_t \sim N(0, 1)$ with probability α and $y_t \sim N(10, 1)$ with probability $(1 - \alpha)$, for $\alpha = 2/3$, $t = 1, \dots, \tau$ and for $\alpha = 1/3$, $t = \tau + 1, \dots, n$ after the change.

In the OU scenario, the data is generated by simulating a discrete Ornstein-Uhlenbeck process (also known as mean-reverting random walk) plus additional Gaussian noise. This is $y_t = \nu_t + \epsilon_t$, where $\epsilon_t \sim N(0, \sigma_\epsilon^2)$ is a white noise process and ν_t is an autoregressive process defined by the recursion $\nu_t = \nu_{t-1} - \theta f_{t-1} - \theta \nu_{t-1} + \sigma_\nu w_{t-1}$, where $\theta = 0.1$, $w_t \sim N(0, \sigma_w^2)$. The sequence f_t is added to encode a shift in the mean of the Ornstein-Uhlenbeck process: for $t = 1, \dots, \tau$, $f_t = 0$, and for $t = \tau + 1, \dots, n$, $f_t = -10$.

The sinusoidal scenario is generated by sampling $y_t \sim N(\mu_t, 1)$ for $t = 1, \dots, n$. The mean μ_t is given by the function $\mu_t = \sin(\pi f t)$, where $f = 0.2$ is the frequency of the sinusoidal function. From the changepoint τ , the mean starts to decrease exponentially with rate $\lambda = 0.005$, i.e. for $t = \tau + 1, \dots, n$, $\mu_t = A \sin(\pi f t) \exp(-\lambda t)$.

Finally, the change in tails scenario is generated by sampling $y_t \sim t_\nu$ for $t = 1, \dots, n$, where t_ν is a Student's t-distribution with ν degrees of freedom. After the changepoint, for $t = \tau + 1, \dots, n$, we randomly selected a 5^{th} of the observations and increase by a random amount drawn from a Poisson distribution with mean 10.

C. Additional Simulation Results

In Section 5 we have commented on the choice of M , the number of quantiles. As for the quantile points, when those cannot be provided directly by the practitioners, those needs to be estimated through a probation period. We therefore performed a simulation study to see the effect of the probation period (amount of training observations) for the quantile estimation on the performances of NP-FOCuS. As in the main study, we present detection delay and false positive rate in Table II(a) and II(b). We found that, in general, even as little as 100 observations are enough for estimating the quantiles. In many cases the length of the probation period has marginal impact on the performances. The major differences are seen on the sinusoidal scenario, where 25 and 50 observations are not enough to capture the full domain of the distribution, resulting in significantly more false positives or slower detections.

To see how the methods deal with an earlier change, we replicate the study of Section IV-A, but we shift the change at time $\tau = 800$. Results are summarised in Fig. 9. We denote how there is a significant drop in performances in NEWMA

TABLE II
AVERAGE DETECTION DELAY AND FALSE POSITIVE RATE FOR ALL SCENARIOS OF OUR SIMULATION STUDY OVER A RANGE OF DIFFERENT PROBATION PERIODS

scenario	25	50	100	150	200
OUmean	102.97	94.29	87.99	78.46	80.71
cauchy	46.19	46.45	33.98	34.47	43.82
gauss	29.57	28.3	22.26	27.4	26.58
multim	71.6	61.24	44.86	55.9	60.88
sinusoidal	154.68	252.3	165.8	140.92	130.61
tails	36.55	55.85	46.97	49.33	44.4

(a) Detection delay

scenario	25	50	100	150	200
OUmean	0	0	0	0.01	0.01
cauchy	0	0	0.01	0.01	0
gauss	0	0	0.01	0	0
multim	0.01	0.01	0.03	0.02	0.01
sinusoidal	0.08	0.03	0.03	0.04	0.05
tails	0.02	0	0	0	0

(b) False positive rate

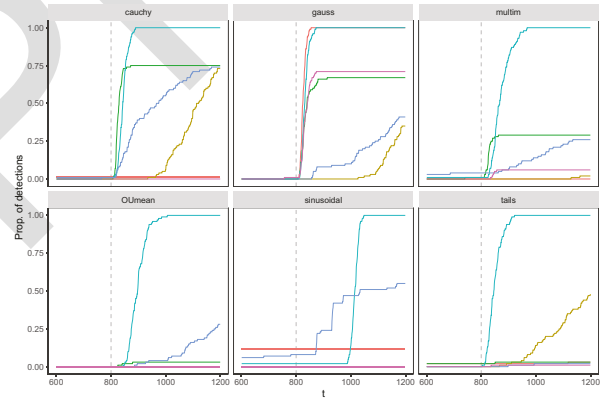


Fig. 9. Proportions of changes detected within t observations following the change in six different change scenarios. The change is denoted by the vertical dotted line at $t = 800$.

across all simulations scenarios. This is particularly evident in the Cauchy and multimodal scenario. The reduction in detection power is attributable to a smaller number of random features, in order to stabilise the statistics faster and achieve shorter burn-in periods. This result is in line with what observed in Section VI-B of [25].

ACKNOWLEDGMENT

The authors would like to thank Trevor Burbridge at BT for the helpful conversations that helped motivate the presented procedure. We are grateful to the Associate Editor and reviewers for their thoughtful comments and suggestions.

REFERENCES

- [1] J. S. Abramson, "Some minorants and majorants of random walks and Lévy processes," Ph.D. thesis, UC Berkeley, 2012.
- [2] J. Alvarez-Montoya, A. Carvajal-Castrillón, and J. Sierra-Pérez, "In-flight and wireless damage detection in a UAV composite wing using fiber optic sensors and strain field pattern recognition," *Mech. Syst. Signal Process.*, vol. 136, 2020, Art. no. 106526.

AQ11

AQ12

AQ13

- [3] E. S. Andersen, "On the fluctuations of sums of random variables II," *Math. Scand.*, pp. 195–223, 1955.
- [4] E. Austin, G. Romano, I. A. Eckley, and P. Fearnhead, "Online non-parametric changepoint detection with application to monitoring operational performance of network devices," *Comput. Statist. Data Anal.*, vol. 177, 2023, Art. no. 107551.
- [5] M. Basseville and I. V. Nikiforov, *Detection of Abrupt Changes: Theory and Application*, vol. 104. Englewood Cliffs, NJ, USA: Prentice Hall, 1993.
- [6] H. Chen, "Sequential change-point detection based on nearest neighbors," *Ann. Statist.*, vol. 47, no. 3, pp. 1381–1407, 2019.
- [7] H. Chen and L. Chu, "Graph-based change-point analysis," *Annu. Rev. Statist. Appl.*, vol. 10, pp. 475–499, 2023.
- [8] Y. Chen, T. Wang, and R. J. Samworth, "High-dimensional, multiscale online changepoint detection," 2020, *arXiv:2003.03668*.
- [9] Y. Chen, T. Wang, and R. J. Samworth, "High-dimensional, multiscale online changepoint detection," *J. Roy. Statist. Soc. Ser. B: Statist. Methodol.*, vol. 84, no. 1, pp. 234–266, 2022.
- [10] H. Cho and P. Fryzlewicz, "Multiple change point detection under serial dependence: Wild energy maximisation and gappy Schwarz criterion," 2020, *arXiv:2011.13884*.
- [11] H. Cho and C. Kirch, "Data segmentation algorithms: Univariate mean change and beyond," *Econometrics Statist.*, 2021.
- [12] L. Chu and H. Chen, "Sequential change-point detection for high-dimensional and non-Euclidean data," *IEEE Trans. Signal Process.*, vol. 70, pp. 4498–4511, 2022.
- [13] P. Fearnhead and P. Fryzlewicz, "Detecting a single change-point," 2022, *arXiv:2210.07066*.
- [14] P. Fearnhead and G. Rigai, "Relating and comparing methods for detecting changes in mean," *Stat.*, vol. 9, no. 1, 2020.
- [15] T. Flynn and S. Yoo, "Change detection with the kernel cumulative sum algorithm," in *Proc. IEEE 58th Conf. Decis. Control (CDC)*, Piscataway, NJ, USA: IEEE, 2019, pp. 6092–6099.
- [16] P. Fridman, "A method of detecting radio transients," *Mon. Not. Roy. Astron. Soc.*, vol. 409, no. 2, pp. 808–820, 2010.
- [17] P. Fryzlewicz and S. S. Rao, "Multiple-change-point detection for autoregressive conditional heteroscedastic processes," *J. Roy. Statist. Soc. Ser. B: Statist. Methodol.*, pp. 903–924, 2014.
- [18] F. Fuschino et al., "Hermes: An ultra-wide band x and gamma-ray transient monitor on board a nano-satellite constellation," *Nucl. Instrum. Methods Phys. Res. Sect. A*, vol. 936, pp. 199–203, 2019.
- [19] K. L. Hallgren, N. A. Heard, and N. M. Adams, "Changepoint detection in non-exchangeable data," 2021, *arXiv:2111.05054*.
- [20] D. M. Hawkins and Q. Deng, "A nonparametric change-point control chart," *J. Qual. Technol.*, vol. 42, no. 2, pp. 165–173, 2010.
- [21] K. Haynes, I. A. Eckley, and P. Fearnhead, "Computationally efficient changepoint detection for a range of penalties," *J. Comput. Graphical Statist.*, vol. 26, no. 1, pp. 134–143, 2017.
- [22] D. Henry, S. Simani, and R. J. Patton, "Fault detection and diagnosis for aeronautic and aerospace missions," in *Fault Tolerant Flight Control*. Springer, 2010, pp. 91–128.
- [23] S. Huang, Z. Kong, and W. Huang, "High-dimensional process monitoring and change point detection using embedding distributions in reproducing kernel Hilbert space," *IIE Trans.*, vol. 46, no. 10, pp. 999–1016, 2014.
- [24] S. W. Jewell, T. D. Hocking, P. Fearnhead, and D. M. Witten, "Fast non-convex deconvolution of calcium imaging data," *Biostatistics*, vol. 21, no. 4, pp. 709–726, 2020.
- [25] N. Keriven, D. Garreau, and I. Poli, "NEWMA: A new method for scalable model-free online change-point detection," *IEEE Trans. Signal Process.*, vol. 68, pp. 3515–3528, 2020.
- [26] S. Kovács, H. Li, P. Bühlmann, and A. Munk, "Seeded binary segmentation: A general methodology for fast and optimal change point detection," 2020, *arXiv:2002.06633*.
- [27] M. N. Kurt, Y. Yimaz, and X. Wang, "Real-time nonparametric anomaly detection in high-dimensional settings," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 43, no. 7, pp. 2463–2479, Jul. 2020.
- [28] D. S. Matteson and N. A. James, "A nonparametric approach for multiple change point analysis of multivariate data," *J. Amer. Statist. Assoc.*, vol. 109, no. 505, pp. 334–345, 2014.
- [29] Y. Mei, "Efficient scalable schemes for monitoring a large number of data streams," *Biometrika*, vol. 97, no. 2, pp. 419–433, 2010.
- [30] P. Nicolas, A. Leduc, S. Robin, S. Rasmussen, H. Jarmer, and P. Bessières, "Transcriptional landscape estimation from tiling array data using a model of signal shift and drift," *Bioinformatics*, vol. 25, no. 18, pp. 2341–2347, 2009.
- [31] E. S. Page, "Continuous inspection schemes," *Biometrika*, vol. 41, no. 1/2, pp. 100–115, 1954.
- [32] A. N. Pettitt, "A non-parametric approach to the change-point problem," *J. Roy. Statist. Soc. Ser. C (Appl. Statist.)*, vol. 28, no. 2, pp. 126–135, 1979.
- [33] A. Pouliezios and G. S. Stavrakakis, *Real Time Fault Monitoring of Industrial Processes*, vol. 12. Springer Science & Business Media, 2013.
- [34] G. Romano, I. A. Eckley, P. Fearnhead, and G. Rigai, "Fast online changepoint detection via functional pruning CUSUM statistics," *J. Mach. Learn. Res.*, vol. 24, no. 81, pp. 1–36, 2023.
- [35] G. Romano, G. Rigai, V. Runge, and P. Fearnhead, "Detecting abrupt changes in the presence of local fluctuations and autocorrelated noise," *J. Amer. Statist. Assoc.*, pp. 1–16, 2021.
- [36] G. J. Ross, "Parametric and nonparametric sequential change detection in R: The cpm package," *J. Statist. Softw.*, vol. 66, pp. 1–20, 2015.
- [37] G. J. Ross, "Nonparametric detection of multiple location-scale change points via wild binary segmentation," 2021, *arXiv:2107.01742*.
- [38] G. J. Ross and N. M. Adams, "Two nonparametric control charts for detecting arbitrary distribution changes," *J. Qual. Technol.*, vol. 44, no. 2, pp. 102–116, 2012.
- [39] A. J. Scott and M. Knott, "A cluster analysis method for grouping means in the analysis of variance," *Biometrics*, vol. 30, no. 3, pp. 507–512, 1974.
- [40] J. Shin, A. Ramdas, and A. Rinaldo, "E-detectors: A nonparametric framework for online changepoint detection," 2022, *arXiv:2203.03532*.
- [41] A. G. Tartakovsky, B. L. Rozovskii, and K. Shah, "A nonparametric multichart CUSUM test for rapid intrusion detection," in *Proc. Joint Statist. Meet.*, vol. 7, 2005, p. 11.
- [42] H. Wang and Y. Xie, "Sequential change-point detection: Computation versus statistical performance," 2022 *arXiv:2210.05181*.
- [43] K. Ward, G. Romano, I. Eckley, and P. Fearnhead, "A constant-per-iteration likelihood ratio test for online changepoint detection for exponential family models," 2023 *arXiv:2302.04743*.
- [44] S. Wei and Y. Xie, "Online kernel CUSUM for change-point detection," 2022, *arXiv:2211.15070*.
- [45] L. Xie, G. V. Moustakides, and Y. Xie, "Window-limited CUSUM for sequential change detection," *IEEE Trans. Inf. Theory*, vol. 69, no. 9, pp. 5990–6005, Sep. 2023.
- [46] Y. Yu, O. M. Padilla, D. Wang, and A. Rinaldo, "A note on online change point detection," 2020, *arXiv:2006.03283*.
- [47] C. Zou, G. Yin, L. Feng, and Z. Wang, "Nonparametric maximum likelihood approach to multiple change-point problems," *Ann. Statist.*, vol. 42, no. 3, pp. 970–1002, 2014.